

Anomaly detection using Gaussian Mixture Model

Bùi Quang Long

February 2020

Contents

1	Expectation maximization	1
1.1	Proof of equation (2)	1
1.2	Formula for n-D data	2
1.2.1	E step	2
1.2.2	M step	3
2	Anomaly detection	3
2.1	Optimal number of clusters (K)	3
2.2	The likelihood of an abnormal sample	4
2.3	Threshold	4
3	Implementation and Explanation	4
3.1	Visualization	4
3.2	GMM Implementation	5
3.2.1	Data preparation	5
3.2.2	Model building	5
3.2.3	Result and Comparison with sklearn	6
3.2.4	Anomaly detection	6

1 Expectation maximization

1.1 Proof of equation (2)

$$b_k = \frac{f(x | \mu_k, \sigma_k^2) \phi_k}{\sum_{k=1}^K f(x | \mu_k, \sigma_k^2) \phi_k}$$

First, we define:

- $p(x)$: the likelihood of an random observation equals to x .
- ϕ_k : the prior beliefs that an example was drawn from the k_{th} (or k^{th} Gaussian distribution).

- $p(z_k)$: the likelihood of an random observation belongs to the k^{th} cluster. (*prior probability*). Here,

$$p(z_k) = \phi_k$$

- $p(x | z_k) = f(x | \mu_k, \sigma_k^2)$: the likelihood of an observation x using the estimated parameters of the k^{th} cluster. (*prior probability*)
- $p(z_k | x)$ (or b_k in equation (2)): the likelihood of a given observation x belongs to the k_{th} cluster. (*posterior probability*). Thus,

$$p(x) = \sum_{j=1}^K p(z_j, x)$$

- $p(z_k, x) = p(x, z_k)$: the likelihood of a random observation satisfies both 2 conditions: equals to x and belongs to the k_{th} cluster. Applying Bayes' Rule,

$$p(z_k, x) = p(x | z_k)p(z_k)$$

Then, we obtain:

$$p(z_k | x) = \frac{p(z_k, x)}{p(x)}$$

$$p(z_k | x) = \frac{p(z_k, x)}{\sum_{j=1}^K p(z_j, x)}$$

$$p(z_k | x) = \frac{p(x | z_k)p(z_k)}{\sum_{j=1}^K p(x | z_j)p(z_j)}$$

$$p(z_k | x) = \frac{f(x | \mu_k, \sigma_k^2)\phi_k}{\sum_{j=1}^K f(x | \mu_j, \sigma_j^2)\phi_j}$$

$$b_k = \frac{f(x | \mu_k, \sigma_k^2)\phi_k}{\sum_{j=1}^K f(x | \mu_j, \sigma_j^2)\phi_j}$$

1.2 Formula for n-D data

1.2.1 E step

- $f(x | \mu, \Sigma)$: The likelihood of an observation x using the estimated parameters of the k^{th} cluster. Here,
 - μ is a N-dimensional vector denoting the mean of the distribution.
 - Σ is the NxN symmetric covariance matrix and positive definite.

$$f(x | \mu, \Sigma) = \frac{\exp(-\frac{1}{2}(x - \mu)^T \Sigma^{-1}(x - \mu))}{\sqrt{(2\pi)^k |\Sigma|}}$$

- b_k : The likelihood of a given observation belongs to the k_{th} cluster.

$$b_k = \frac{f(x | \mu_k, \Sigma_k) \phi_k}{\sum_{j=1}^K f(x | \mu_j, \Sigma_j) \phi_j}$$

For further calculation, we define b_{nk} equals to the likelihood of the n_{th} observation belongs to the k_{th} cluster.

$$b_{nk} = \frac{f(x_n | \mu_k, \Sigma_k) \phi_k}{\sum_{j=1}^K f(x_n | \mu_j, \Sigma_j) \phi_j}$$

1.2.2 M step

Here we define N is number of observation in our data set. After E step, we recalculate the following parameters:

$$\begin{aligned}\mu_k^{new} &= \frac{1}{N_k} \sum_{n=1}^N b_{nk} x_n \\ \Sigma_k^{new} &= \frac{1}{N_k} \sum_{n=1}^N b_{nk} (x_n - \mu_k^{new})(x_n - \mu_k^{new})^T \\ \phi_k^{new} &= \frac{N_k}{N}\end{aligned}$$

where

$$N_k = \sum_{n=1}^N b_{nk}$$

2 Anomaly detection

2.1 Optimal number of clusters (K)

To define a good number of clusters, we can use the Bayesian information criterion (BIC) or the Akaike information criterion (AIC). Here is my quick comparison between these two criteria:

- AIC is more suitable for unstable data (small, large noise levels, ...).
- AIC tries to select the model that most adequately describes an unknown, high dimensional reality. On the contrary, BIC tries to find the TRUE model among the set of candidates.
- BIC penalizes models more for free parameters than does AIC.

Hence our training data has more than normal 1500 samples, I use the BIC and get the result: The model works best at K around 5 to 20. (I choose 7)




Figure 1: Using BIC to define number of clusters

2.2 The likelihood of an abnormal sample

The likelihood of an sample x :

$$p(x) = \sum_{k=1}^K f(x | \mu_k, \Sigma_k) \phi_k$$

2.3 Threshold

I come up with a simple threshold formula based on normal samples: find the min likelihood among normal samples.

$$threshold = \min(p(x))$$

3 Implementation and Explanation

3.1 Visualization

The data has 21 dimensions for each sample so the first thing I come up with is dimensionality reduction, which led me to Principal component analysis (PCA). With 3 dimensions, I get:




Figure 2: 3D illustration

It looks like lots of spatial information can't be illustrated. Thus I haven't recognized the separation clearly. More detail will be illustrated in Result Section.

3.2 GMM Implementation

3.2.1 Data preparation

- After several attempts, I was stuck with overflow error. The covariance matrices became too small and overflow the float type. Thus I use PCA to reduce the numbers of features to 13 and still maintain 95% of original information.
- From the data set of 1831 samples, I also split the data into 2 parts: normal and suspect part. Now my samples for training has the shape of (1655, 13)

3.2.2 Model building

- I apply the Expectation-Maximization algorithm. The goal is to maximize the Log-Likelihood function:

$$\ln p(X | \mu, \Sigma, \phi) = \sum_{n=1}^N \ln \left(\sum_{k=1}^k \phi_k f(x_n | \mu_k, \Sigma_k) \right)$$

```

11 # todo part1: Load data
12 data = loadmat('cardio.mat')
13 X = np.array(data['X'])
14 y = data['y']
15 X_train = X[np.where(y == 0)[0]]
16 X_test = X[np.where(y == 1)[0]]
17 # Dimensionality reduction: maintain 95% of original information. The result data shape is (1655, 13).
18 # Reason: To avoid overflow calculation
19 pca = PCA(0.95, whiten=True, random_state=0)
20 X_train = pca.fit_transform(X_train)

```

Figure 3: Data preparation

- I set the convergence criterion to 0.1, the loop will stop when the difference between the Log-Likelihood of this loop and previous loop less than 0.1:

$$curLogLikelihood - preLogLikelihood < 0.1$$

3.2.3 Result and Comparison with sklearn

After 71 iterations, my GMM converged and release the score approximately the score found by sklearn.

```

135 #todo part3: Compare with sklearn
136 from sklearn.mixture import GaussianMixture
137 gmm = GaussianMixture(n_components=nCluster, covariance_type='full', random_state=0)
138 gmm.fit(X_train)
139 print("Score found by my code: ", curLogLikelihood / nData)
140 print("Score found by sklearn: ", gmm.score(X_train))

```

Figure 4: GMM using sklearn

```

70
-10.555433316444343
71
-10.555399186107817
Score found by my code: -10.555399186107817
Score found by sklearn: -10.478886588135985

```

Figure 5: Compare

3.2.4 Anomaly detection

With my threshold formula, I detected 30 anomlies, which will be marked as black points in 3D illustration.




Figure 6: Anomalies distribution




Figure 7: Anomalies and normal samples